

CS 445: Computational Photography

Programming Project #3: Gradient Domain Fusion

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import os
from random import random
import time
import scipy
import scipy.sparse.linalg

import utils
```

Part 1 Toy Problem (20 pts)

```
def toy_reconstruct(img):
    """
    The implementation for gradient domain processing is not
    complicated, but it is easy to make a mistake, so let's start with a
    toy example. Reconstruct this image from its gradient values, plus one
    pixel intensity. Denote the intensity of the source image at (x, y) as
     $s(x,y)$  and the value to solve for as  $v(x,y)$ . For each pixel, then, we
    have two objectives:
    1. minimize  $(v(x+1,y)-v(x,y) - (s(x+1,y)-s(x,y)))^2$ 
    2. minimize  $(v(x,y+1)-v(x,y) - (s(x,y+1)-s(x,y)))^2$ 
    Note that these could be solved while adding any constant value to
     $v$ , so we will add one more objective:
    3. minimize  $(v(1,1)-s(1,1))^2$ 

    :param toy_img: numpy.ndarray
    """

    # Initialization
    H, W = img.shape
    im2var = np.arange(H*W).reshape(H, W)

    total = 2*(H-1)*(W-1)+W+H-1
    A = np.zeros((total, H*W))
    b = np.zeros(total)
    e = 0

    # object 3
    A[e][im2var[0][0]] = 1
    b[e] = img[0][0]
    e += 1
```

```

for x in range(W):
    for y in range(H):

        # object 1
        if x < W - 1:
            A[e][im2var[y][x+1]] = 1
            A[e][im2var[y][x]] = -1
            b[e] = img[y][x+1] - img[y][x]
            e += 1



        # object 2
        if y < H - 1:
            A[e][im2var[y+1][x]] = 1
            A[e][im2var[y][x]] = -1
            b[e] = img[y+1][x] - img[y][x]
            e += 1

v = scipy.sparse.linalg.lsqr(A, b, atol=1e-14, btol=1e-14)[0]
return v.reshape(H, W)

toy_img = cv2.cvtColor(cv2.imread('samples/toy_problem.png'),
cv2.COLOR_BGR2GRAY).astype('double') / 255.0
plt.imshow(toy_img, cmap="gray")
plt.show()

im_out = toy_reconstruct(toy_img)
plt.imshow(im_out, cmap="gray")
plt.show()
print("Max error is: ", np.sqrt(((im_out - toy_img)**2).max()))

```



Max error is: 7.420508651989621e-12

Part 2 Poisson Blending (50 pts)

Preparation

```
background_img = cv2.cvtColor(cv2.imread('samples/universe.jpg'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
# background_img = cv2.cvtColor(cv2.imread('samples/horizon.jpg'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
plt.figure()
plt.imshow(background_img)
plt.show()
object_img = cv2.cvtColor(cv2.imread('samples/earth.png'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
# object_img = cv2.cvtColor(cv2.imread('samples/dolphin2.png'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
# object_img = cv2.cvtColor(cv2.imread('samples/dolphin.png'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
plt.imshow(object_img)
plt.show()

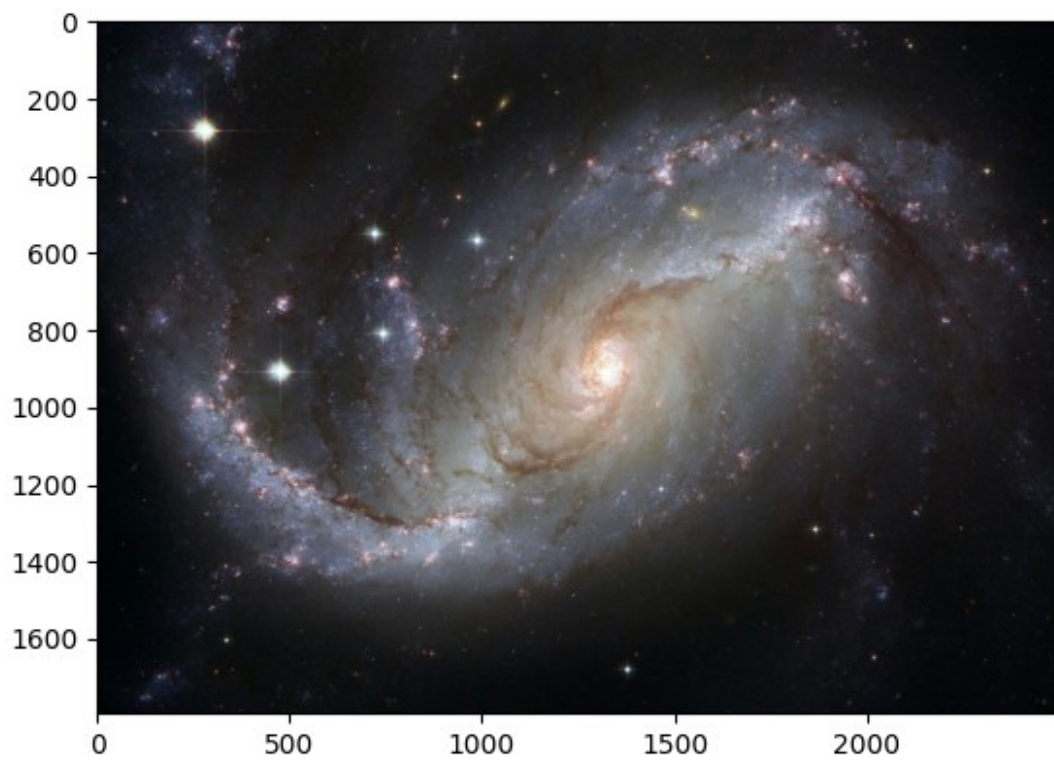
use_interface = False # set to true if you want to use the interface
to choose points (might not work in Colab)
if not use_interface:
    # result 1
    xs = (174, 110, 67, 29, 50, 102, 189, 329, 450, 482, 495, 470,
424, 342, 247)
    ys = (35, 66, 106, 200, 290, 343, 383, 411, 383, 332, 232, 144,
78, 29, 15)
    bottom_center = (1500, 1600) # (x,y)



    # result 2
    # xs = (62, 36, 8, 10, 41, 57, 84, 116, 175, 186, 210, 213, 222,
188, 160, 163, 150, 117, 89)
    # ys = (7, 21, 30, 43, 56, 77, 81, 80, 99, 126, 130, 108, 77, 74,
58, 38, 19, 6, 4)
    # bottom_center = (150, 530) # (x,y)

    # result 3
    # xs = (60, 46, 27, 19, 4, 4, 25, 53, 60, 81, 97, 128, 130, 118,
113, 104, 101, 107, 127, 114, 124, 140, 136, 117, 89)
    # ys = (8, 21, 47, 59, 72, 85, 97, 130, 151, 164, 176, 175, 166,
161, 143, 115, 98, 79, 78, 51, 38, 25, 9, 5, 2)
    # bottom_center = (700, 530) # (x,y)

    object_mask = utils.get_mask(ys, xs, object_img)
    object_img, object_mask = utils.crop_object_img(object_img,
object_mask)
    bg_ul = utils.upper_left_background_rc(object_mask, bottom_center)
    combined_img = utils.get_combined_img(background_img, object_img,
object_mask, bg_ul)
```

```
plt.imshow(combined_img)  
# plt.imsave("./results/part2_simple_result1.png", combined_img)
```





Poisson Blending

```
def poisson_blend(object_img, object_mask, bg_img, bg_ul):  
    """  
    Returns a Poisson blended image with masked object_img over the  
    bg_img at position specified by bg_ul.  
    Can be implemented to operate on a single channel or multiple  
    channels  
    :param object_img: the image containing the foreground object  
    :param object_mask: the mask of the foreground object in  
    object_img  
    :param background_img: the background image  
    :param bg_ul: position (row, col) in background image  
    corresponding to (0,0) of object_img  
    """  
    # get all nonzeros indices  
    rows, cols = np.nonzero(object_mask)  
  
    # construct im2var  
    im2var = np.full(object_mask.shape, -1, dtype=int)  
    for i, (r, c) in enumerate(zip(rows, cols)):  
        im2var[r][c] = i  
  
    data, row_ind, col_ind = [], [], []  
    b = np.zeros(4*len(rows))  
    e = 0  
    increments = [(-1,0), (1,0), (0,-1), (0,1)]  
  
    for r, c in zip(rows, cols):  
        for inc in increments:  
            nr, nc = r+inc[0], c+inc[1]  
  
            row_ind.append(e)  
            col_ind.append(im2var[r][c])  
            data.append(1)  
  
            if object_mask[nr][nc]:  
                row_ind.append(e)  
                col_ind.append(im2var[nr][nc])  
                data.append(-1)  
                b[e] = object_img[r][c] - object_img[nr][nc]  
            else:  
                tnr, tnc = bg_ul[0]+nr-1, bg_ul[1]+nc-1  
                b[e] = bg_img[tnr][tnc] + object_img[r][c] -  
object_img[nr][nc]  
  
                e += 1  
  
    A = scipy.sparse.csr_matrix((data, (row_ind, col_ind)), shape=(e,  
len(rows)))
```



```

v = scipy.sparse.linalg.lsqr(A, b, atol=1e-14, btol=1e-14)[0]
v = np.clip(v, 0.0, 1.0)

# paste to background image
for i, (r, c) in enumerate(zip(rows, cols)):
    tr, tc = bg_ul[0]+r-1, bg_ul[1]+c-1
    bg_img[tr][tc] = v[i]


return bg_img

im_blend = np.zeros(background_img.shape)
for b in np.arange(3):
    im_blend[:, :, b] = poisson_blend(object_img[:, :, b], object_mask,
background_img[:, :, b].copy(), bg_ul)

plt.figure(figsize=(8,8))
plt.imshow(im_blend)

<matplotlib.image.AxesImage at 0x762e2bb25d80>

```



```

# plt.imshow("./results/part2_result1.png", im_blend)

```



Part 3 Mixed Gradients (20 pts)

Preparation

```
background_img = cv2.cvtColor(cv2.imread('samples/wall.jpg'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
plt.figure()
plt.imshow(background_img)
plt.show()
object_img = cv2.cvtColor(cv2.imread('samples/writing.png'),
cv2.COLOR_BGR2RGB).astype('double') / 255.0
plt.imshow(object_img)
plt.show()

use_interface = False # set to true if you want to use the interface
to choose points (might not work in Colab)
if not use_interface:
    # result 1
    xs = (135, 80, 42, 33, 16, 21, 35, 86, 129, 146, 196, 227, 331,
363, 405, 415, 436, 419, 371, 333, 230)
    ys = (7, 77, 113, 180, 313, 410, 454, 464, 407, 348, 346, 379,
374, 434, 401, 336, 260, 167, 136, 37, 11)
    bottom_center = (500, 530) # (x,y)

    object_mask = utils.get_mask(ys, xs, object_img)
    object_img, object_mask = utils.crop_object_img(object_img,
object_mask)
    bg_ul = utils.upper_left_background_rc(object_mask, bottom_center)
    combined_img = utils.get_combined_img(background_img, object_img,
object_mask, bg_ul)
    plt.imshow(combined_img)
    plt.imsave("./results/part3_simple_result.png", combined_img)
```





Mixed Gradients Blending

```
def mixed_blend(object_img, object_mask, bg_img, bg_ul):
    """
    Returns a mixed gradient blended image with masked object_img over
    the bg_img at position specified by bg_ul.
    Can be implemented to operate on a single channel or multiple
    channels
    :param object_img: the image containing the foreground object
    :param object_mask: the mask of the foreground object in
    object_img
    :param background_img: the background image
    :param bg_ul: position (row, col) in background image
    corresponding to (0,0) of object_img
    """
    def get_grad(r, c, nr, nc):
        s_grad = object_img[r][c] - object_img[nr][nc]

        tr, tc = bg_ul[0]+r-1, bg_ul[1]+c-1
        tnr, tnc = bg_ul[0]+nr-1, bg_ul[1]+nc-1
        t_grad = bg_img[tr][tc] - bg_img[tnr][tnc]

        return s_grad if abs(s_grad) > abs(t_grad) else t_grad

    # get all nonzeros indices
    rows, cols = np.nonzero(object_mask)

    # construct im2var
    im2var = np.full(object_mask.shape, -1, dtype=int)
    for i, (r, c) in enumerate(zip(rows, cols)):
```

```

        im2var[r][c] = i

    data, row_ind, col_ind = [], [], []
    b = np.zeros(4*len(rows))
    e = 0
    increments = [(-1,0), (1,0), (0,-1), (0,1)]

    for r, c in zip(rows, cols):
        for inc in increments:
            nr, nc = r+inc[0], c+inc[1]

            row_ind.append(e)
            col_ind.append(im2var[r][c])
            data.append(1)

            grad = get_grad(r, c, nr, nc)

            if object_mask[nr][nc]:
                row_ind.append(e)
                col_ind.append(im2var[nr][nc])
                data.append(-1)
                b[e] = grad
            else:
                tnr, tnc = bg_ul[0]+nr-1, bg_ul[1]+nc-1
                b[e] = bg_img[tnr][tnc] + grad

            e += 1

    A = scipy.sparse.csr_matrix((data, (row_ind, col_ind)), shape=(e,
len(rows)))
    v = scipy.sparse.linalg.lsqr(A, b, atol=1e-14, btol=1e-14)[0]
    v = np.clip(v, 0.0, 1.0)

    # paste to background image
    for i, (r, c) in enumerate(zip(rows, cols)):
        tr, tc = bg_ul[0]+r-1, bg_ul[1]+c-1
        bg_img[tr][tc] = v[i]

    return bg_img

im_mix = np.zeros(background_img.shape)
for b in np.arange(3):
    im_mix[:, :, b] = mixed_blend(object_img[:, :, b], object_mask,
background_img[:, :, b].copy(), bg_ul)

plt.figure(figsize=(8,8))
plt.imshow(im_mix)

<matplotlib.image.AxesImage at 0x762e85ebda0>

```




```
plt.imsave("./results/part3_result.png", im_mix)
```

Bells & Whistles (Extra Points)

Color2Gray (20 pts)

```
def color2gray(img):  
    pass
```

Laplacian pyramid blending (20 pts)

```
def laplacian_blend(object_img, object_mask, bg_img, bg_ul):  
    # feel free to change input parameters  
    pass
```

More gradient domain processing (up to 20 pts)